

The Flower Classification Challenge

This course focuses on the principles underlying deep learning. Therefore, the exercises typically ask you to re-implement known algorithms and consider somewhat simplified settings. Beyond knowledge of these foundations, a deep learning practitioner must also build up experience using the many powerful tools available in the real world. We organize this competition to give you the opportunity to start doing so.

Participation is optional and your submissions will not be graded.

To avoid extra organisational overhead, you will compete in the same teams as for the exercises. The winners will be announced during the last flipped classroom session (10.02.2021) and will receive nothing but eternal glory and a spot on [the course webpage](#). To mimick a real-world setting, we give you a lot of freedom in how you tackle this challenge. However, with this freedom comes responsibility. So please play fair and make it easy for us to evaluate your submission. In doubt, please [contact us](#) or create a thread on the [ILIAS forum](#).

To get access to the code, use the following link: <https://classroom.github.com/g/guOphWsx>

The Challenge:

In this challenge, you are to train a model to perform class prediction on a flower dataset. Specifically, given an image of a flower as input, your model must predict to which of 10 possible classes it belongs.



An example of the flower dataset

To encourage everybody to participate, we have two tracks in this competition:

- Fast networks, < 100k learnable parameters.
- Large networks, < 25M learnable parameters.

These two tracks shall be considered as separate competitions with separate winners. You may compete in either (or both) tracks.

You are allowed to implement any architecture, and manually or automatically tune the hyperparameters of the model. The following are just a few of the things you can experiment with:

- Learning rate and its optional scheduler
- Different optimizers and their hyperparameters
- Specialized model architectures
- Activation functions
- Regularization
- Data preprocessing
- Cross-validation to ensure you don't overfit to the validation set

Also, you don't have to shy away from more advanced techniques like warm-starting your model with weights from other pre-trained models. You are allowed to use all the scripts and tools you already know from the exercises. However, you are not limited to them.

Here are a few important topics concerning the competition:

- **Evaluation:**
 - The final performance will be measured in terms of classification accuracy of your model on an **unseen** test set.
 - To get a taste of how we shall evaluate your model, try running `evaluate_model.py`. The script loads your saved model from `models` folder and evaluates its accuracy on validation data loaded from the `dataset/test` folder.
 - To evaluate your model, we will populate this folder `dataset/test` with the unseen test data.
- **Hardware:**
 - The fast networks can be trained on a CPU in a few minutes. The larger networks might require GPUs.
 - You may use any kind of hardware that is available to you. For example, [Google Colab](#) repeatedly offers a VM with a GPU for at most 12 hours at a time for free. [Amazon SageMaker](#) offers quite some resources for free if you are a first-time customer.
 - You are also free to use the pool computers. You can use ssh to log in remotely to these computers as follows: `ssh yourpoolaccount@login.informatik.uni-freiburg.de`
From there, you can log in to pool computers with GPUs using `ssh tfpoolxx` where `xx = {20...64}`. Only a few machines may be up and running at a given moment. You can check if a host is up using this script, which returns 0 if the host is available:

```
for i in {20..64}; do nc -w 1 -z "tfpool${i}" 22 > /dev/null; echo "tfpool${i}: ${?}"; done
```
- **Implementation Constraints:**
 - Your model should be written using PyTorch.
 - Do not modify the code inside `src/eval`. Your model will be evaluated using this code.
 - Other than that, you are free to extend/modify the baseline code given to you or write your own code from scratch.
- **Github Repository Constraints:**
 - **Keep your repository under 150 MB.**
 - Github does not allow files larger than 100MB to be tracked in repositories. **So if you are competing in the Large Networks track, make sure your trained model is under this limit.**
 - You are allowed to push only one trained model per track.
 - Your code has to be in the `master` branch before the deadline.
 - There is no limit on the number of pushes to the `master` branch.

As a starting point, we provide you with the following:

- `dataset` folder containing the original flower classification dataset and train/validation/test split.¹ We use only 10 out of the 17 classes provided by the dataset, so that is what you are given.
- `data_augmentations.py` containing two sample data augmentation pipelines.
- `training.py` containing code to train your model.
- `eval/evaluate.py` containing the code we shall use to evaluate your model. **Do not edit this file.**
- `main.py` containing the code to load the data, train, evaluate, and, optionally, save the model. You can run it from the root directory using `python -m src.main`.
- `cnn.py` containing `SampleModel`, as an example of what a (very basic) convolutional model looks like.
- A sample saved model in `src/models`. This file contains the weights of the `SampleModel` given above trained for 50 epochs using the default pipeline that is provided to you.
- `evaluate_model.py` which the organizers will use to evaluate your model. You can run it from the root directory using `python -m src.evaluate_model`.
- `requirements.txt` which contains the list of libraries required to run the given pipeline. **Please manually install torch and torchvision** (see [assignment_07.pdf](#) for instructions).

¹You are free to consider different splits, use cross-validation, or even use all the data to train your final model(s).

Your submission must include:

- **A fully functional training pipeline.** This must include the code for your model, data augmentations you use and the code for training the model.²
- `requirements.txt` updated with any additional libraries you use.
- **The file(s) with the saved weights of your trained model(s)**, similar to `models/sample_model`, in the `models` folder. The files must be saved as `fast_model` or `large_model`³, depending on the track it belongs to.
- `submission.md` must be populated with the answers to the following questions:
 - The track(s) you are competing in.
 - The number of learnable parameters in your model(s). You can find this out using `torchsummary.summary`.
 - A *brief* description of your approach to the problem.
 - Command to run to train your model(s) from scratch with your hyperparameter settings and data augmentations. For example:

```
python -m src.main --model FastModel1 --epochs 60 --data-augmentation
resize_to_64x64 --use-all-data-to-train
```

You're free to edit `main.py` however you please to make this work (E.g., to accept hyperparameters as options, or hard-code them). You can also add new files, if you wish. **We must be able to train your model by running a single command.**
 - Command to run to evaluate your saved model(s) with your data augmentation pipeline. For example:

```
python -m src.eval_model --model FastModel1 --saved-model-file fast_model
```

Again, you're free to edit `evaluate_model.py` however you want, **but ultimately, evaluation MUST be done by invoking `eval_model(...)` function in `eval/evaluate.py`. We must also be able to evaluate your model by running a single command.**
- `sample_submission.zip` contains a sample of what a submission must look like.

This project is due on 28.01.2021 23:59 CET.

²If you use model warm-starting or other such methods, that should also be included in this pipeline. We must be able to run your code and train your model from scratch, as you did.

³Github has a limit of 100MB for tracked files. Your `large_model` must hence be smaller than 100MB.